

```
In [9]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
```

```
In [10]: df=pd.read_csv("health care diabetes.csv")
```

```
In [11]: df.columns
```

```
Out[11]: Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
              'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
              dtype='object')
```

```
In [12]: df.head()
```

```
Out[12]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age
0	6	148	72	35	0	33.6	0.627	50
1	1	85	66	29	0	26.6	0.351	31
2	8	183	64	0	0	23.3	0.672	32
3	1	89	66	23	94	28.1	0.167	21
4	0	137	40	35	168	43.1	2.288	33



```
In [13]: df.dtypes
```

```
Out[13]: Pregnancies      int64
Glucose      int64
BloodPressure  int64
SkinThickness int64
Insulin      int64
BMI          float64
DiabetesPedigreeFunction float64
Age          int64
Outcome      int64
dtype: object
```

```
In [14]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies           768 non-null   int64
1   Glucose               768 non-null   int64
2   BloodPressure         768 non-null   int64
3   SkinThickness         768 non-null   int64
4   Insulin               768 non-null   int64
5   BMI                   768 non-null   float64
6   DiabetesPedigreeFunction 768 non-null   float64
7   Age                   768 non-null   int64
```

```
8 Outcome
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

In [15]:

```
df.corr()
```

Out[15]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI
Pregnancies	1.000000	0.129459	0.141282	-0.081672	-0.073535	0.017683
Glucose	0.129459	1.000000	0.152590	0.057328	0.331357	0.221071
BloodPressure	0.141282	0.152590	1.000000	0.207371	0.088933	0.281805
SkinThickness	-0.081672	0.057328	0.207371	1.000000	0.436783	0.392573
Insulin	-0.073535	0.331357	0.088933	0.436783	1.000000	0.197859
BMI	0.017683	0.221071	0.281805	0.392573	0.197859	1.000000
DiabetesPedigreeFunction	-0.033523	0.137337	0.041265	0.183928	0.185071	0.140647
Age	0.544341	0.263514	0.239528	-0.113970	-0.042163	0.036242
Outcome	0.221898	0.466581	0.065068	0.074752	0.130548	0.292695

1. Perform descriptive analysis. Understand the variables and their corresponding values.

In [16]:

```
df.describe()
```

Out[16]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	0.332424
std	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	0.332424
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	0.000000
75%	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	0.000000
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	0.671600

In [17]:


```
df.shape
```

Out[17]:

```
(768, 9)
```

In [18]:

```
df.boxplot();
```



```
In [19]: df.isnull().sum()
```

```
Out[19]: Pregnancies      0
Glucose      0
BloodPressure  0
SkinThickness 0
Insulin      0
BMI          0
DiabetesPedigreeFunction 0
Age          0
Outcome      0
dtype: int64
```

```
In [20]: df.columns
```

```
Out[20]: Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
               'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
              dtype='object')
```

a value of zero does not make sense and thus indicates missing

value: Treating the null values

```
In [21]: for i in ['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
                 'BMI']:
           print(i)
           print(df[i].value_counts(normalize=True)[0], '\n\n')
```

```
Glucose
0.006510416666666667
```

```
BloodPressure
0.045572916666666664
```

```
SkinThickness
0.2955729166666667
```

```
Insulin
0.4869791666666667
```

```
BMI
```

0.014322916666666666

```
In [22]: df[df['Insulin']!=0]['Insulin'].median()
```

```
Out[22]: 125.0
```

```
In [23]: for i in ['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
                'BMI']:
            print(i)
            median_value=df[df[i]!=0][i].median()
            print(median_value)
            df[i].replace(0,median_value,inplace=True)
```

```
Glucose
117.0
BloodPressure
72.0
SkinThickness
29.0
Insulin
125.0
BMI
32.3
```

3. There are integer and float data type variables in this dataset. Create a count

(frequency) plot describing the data types and the count of variables.

```
In [24]: df.dtypes.value_counts()
```

```
Out[24]: int64      7
float64      2
dtype: int64
```

```
In [25]: df[df[['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
                'BMI']]==0].count()
```


```
Out[25]: Pregnancies      0
Glucose      0
BloodPressure      0
SkinThickness      0
Insulin      0
BMI      0
DiabetesPedigreeFunction      0
Age      0
Outcome      0
dtype: int64
```

```
In [51]: sns.countplot(df.dtypes.value_counts())
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
<AxesSubplot:ylabel='count'>
```


```
Out[51]:
```



In [26]: `df.dtypes.value_counts()`

Out[26]:
 int64 7
 float64 2
 dtype: int64

In [27]: `plt.figure(figsize=(12,8))`
`df.boxplot();`



In [28]: `df.columns`

Out[28]: `Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'], dtype='object')`

Check the balance of the data by plotting the count of outcomes by their value.

```
In [29]: df.Outcome.value_counts(normalize=True)
```

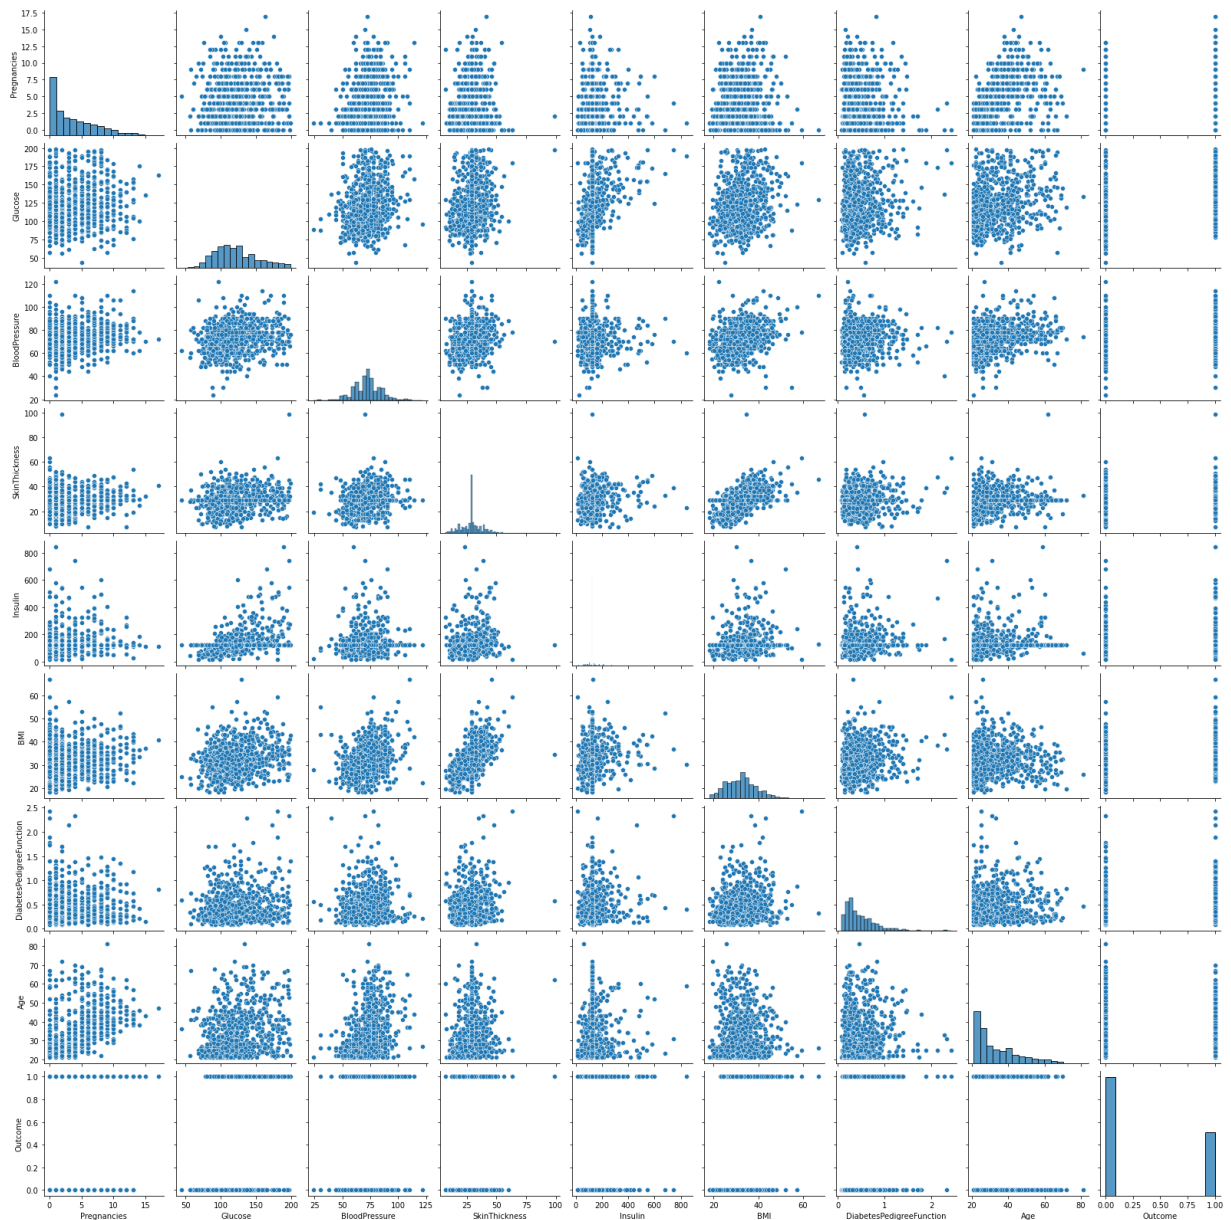
```
Out[29]: 0    0.651042
         1    0.348958
         Name: Outcome, dtype: float64
```

Create scatter charts between the pair of variables to understand the relationships.

Describe your findings.

```
In [30]: sns.pairplot(df)
```

```
Out[30]: <seaborn.axisgrid.PairGrid at 0x1f197d0cf40>
```




```
In [52]: sns.distplot(df)
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)


```
Out[52]: <AxesSubplot:ylabel='Density'>
```



Perform correlation analysis. Visually explore it using a heat map.

```
In [31]: sns.heatmap(df.corr())
```

```
Out[31]: <AxesSubplot:>
```



```
In [32]: df.columns
```

```
Out[32]: Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
              'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
              dtype='object')
```

```
In [33]: x=df.drop('Outcome',axis=1)
          y=df['Outcome']
```

```
In [34]: df['Outcome'].value_counts(normalize=True)
```

```
Out[34]: 0    0.651042
          1    0.348958
          Name: Outcome, dtype: float64
```

```
In [35]: train_x, test_x, train_y, test_y = train_test_split(x, y, test_size=25, random_state=42, stra
```

```
In [36]: train_x.shape
```

```
Out[36]: (743, 8)
```

```
In [37]: test_x.shape
```

```
Out[37]: (25, 8)
```

```
In [38]: test_y.value_counts(normalize=True)
```

```
Out[38]: 0    0.64  
         1    0.36  
         Name: Outcome, dtype: float64
```

```
In [39]: train_y.value_counts(normalize=True)
```

```
Out[39]: 0    0.651413  
         1    0.348587  
         Name: Outcome, dtype: float64
```

```
In [40]: model = KNeighborsClassifier()
```

```
In [41]: model.fit(train_x, train_y)
```

```
Out[41]: KNeighborsClassifier()
```

```
In [42]: model.score(train_x, train_y)
```

```
Out[42]: 0.7981157469717362
```

```
In [43]: model.score(test_x, test_y)
```

```
Out[43]: 0.76
```

Normalizing the Train and test data Using Minmax scalar

```
In [53]: from sklearn.preprocessing import MinMaxScaler
```

```
In [54]: scalar = MinMaxScaler()
```

```
In [56]: scaled_train_x = scalar.fit_transform(train_x)
```

```
In [57]: scaled_test_x = scalar.fit_transform(test_x)
```


KNN classifier model

```
In [59]: from sklearn.neighbors import KNeighborsClassifier
```

```
In [60]: knn_model=KNeighborsClassifier(n_neighbors=20)
```

```
In [62]: knn_model.fit(scaled_train_x,train_y)
```

```
Out[62]: KNeighborsClassifier(n_neighbors=20)
```

```
In [63]: pred_y_test_knn=knn_model.predict(scaled_test_x)
```

```
In [64]: from sklearn.metrics import accuracy_score

print("The accuracy score of the model on test data is : ")
print(accuracy_score(test_y , pred_y_test_knn))
```

The accuracy score of the model on test data is :
0.84

```
In [65]: pred_y_train_knn=knn_model.predict(scaled_train_x)
```

```
In [66]: print("The accuracy score of the model on train data is : ")
print(accuracy_score(train_y , pred_y_train_knn))
```

The accuracy score of the model on train data is :
0.7860026917900403

We could infer that a higher K value would affect the prediction accuracy

```
In [45]: from sklearn.metrics import classification_report,confusion_matrix
```

```
In [46]: y_pred=model.predict(test_x)
```

```
In [47]: print(classification_report(test_y,y_pred))
```

	precision	recall	f1-score	support
0	0.81	0.81	0.81	16
1	0.67	0.67	0.67	9
accuracy			0.76	25
macro avg	0.74	0.74	0.74	25
weighted avg	0.76	0.76	0.76	25

```
In [48]: print(confusion_matrix(test_y,y_pred))
```

```
[[13  3]
 [ 3  6]]
```

Naive Bayes Classifier - Gaussian NB

```
In [68]: from sklearn.naive_bayes import GaussianNB
```

```
In [69]: gnb_model=GaussianNB()
```

```
In [71]: gnb_model.fit(scaled_train_x,train_y)
```

```
Out[71]: GaussianNB()
```

```
In [73]: y_gnb_test=gnb_model.predict(scaled_test_x)
```

```
In [74]: print("The accuracy score of the model on test data is : ")
print(accuracy_score(test_y , y_gnb_test))
```

The accuracy score of the model on test data is :
0.68

```
In [75]: y_gnb_train=gnb_model.predict(scaled_train_x)
```

```
In [78]: print(classification_report(test_y , y_gnb_test))
```

	precision	recall	f1-score	support
0	0.90	0.56	0.69	16
1	0.53	0.89	0.67	9
accuracy			0.68	25
macro avg	0.72	0.73	0.68	25
weighted avg	0.77	0.68	0.68	25

Build SVC Model

```
In [79]: from sklearn.svm import SVC
```

```
In [80]: svc_model=SVC()
```

```
In [82]: svc_model.fit(scaled_train_x,train_y)
```

```
Out[82]: SVC()
```

```
In [83]: svc_model_test_pred=svc_model.predict(scaled_test_x)
```

```
In [85]: svc_model_train_pred=svc_model.predict(scaled_train_x)
```

```
In [87]: print("The accuracy score of the model on test data is : ")
print(accuracy_score(test_y , svc_model_test_pred))
```

The accuracy score of the model on test data is :
0.6

```
In [88]: print("The accuracy score of the model on train data is : ")
print(accuracy_score(train_y , svc_model_train_pred))
```

The accuracy score of the model on train data is :
0.8088829071332436

```
In [90]: print(classification_report(test_y , svc_model_test_pred))
```

	precision	recall	f1-score	support
0	0.67	0.75	0.71	16
1	0.43	0.33	0.38	9
accuracy			0.60	25
macro avg	0.55	0.54	0.54	25
weighted avg	0.58	0.60	0.59	25

```
In [ ]:
```