

Following actions should be performed:

- If for any column(s), the variance is equal to zero, then you need to remove those variable(s).
- Check for null and unique values for test and train sets.
- Apply label encoder.
- Perform dimensionality reduction.
- Predict your test_df values using XGBoost.

In [45]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import warnings
warnings.filterwarnings('ignore')
```

In [46]:

```
df_train=pd.read_csv("train.csv")
df_test=pd.read_csv("test.csv")
```

In [47]:

```
df_train.head()
```

Out[47]:

	ID	y	X0	X1	X2	X3	X4	X5	X6	X8	...	X375	X376	X377	X378	X379	X380	X381
0	0	130.81	k	v	at	a	d	u	j	o	...	0	0	1	0	0	0	0
1	6	88.53	k	t	av	e	d	y	l	o	...	1	0	0	0	0	0	0
2	7	76.26	az	w	n	c	d	x	j	x	...	0	0	0	0	0	0	0
3	9	80.62	az	t	n	f	d	x	l	e	...	0	0	0	0	0	0	0
4	13	78.02	az	v	n	f	d	h	d	n	...	0	0	0	0	0	0	0

5 rows × 378 columns

In [48]:

```
df_train.shape
```

Out[48]:

(4209, 378)


In [49]:

```
df_train.describe()
```

Out[49]:

	ID	y	X10	X11	X12	X13	X14
count	4209.000000	4209.000000	4209.000000	4209.0	4209.000000	4209.000000	4209.000000
mean	4205.960798	100.669318	0.013305	0.0	0.075077	0.057971	0.428130
std	2437.608688	12.679381	0.114590	0.0	0.263547	0.233716	0.494867
min	0.000000	72.110000	0.000000	0.0	0.000000	0.000000	0.000000
25%	2095.000000	90.820000	0.000000	0.0	0.000000	0.000000	0.000000
50%	4220.000000	99.150000	0.000000	0.0	0.000000	0.000000	0.000000
75%	6314.000000	109.010000	0.000000	0.0	0.000000	0.000000	1.000000
max	8417.000000	265.320000	1.000000	0.0	1.000000	1.000000	1.000000

8 rows × 370 columns



In [50]:

```
df_train.columns
```

Out[50]:

```
Index(['ID', 'y', 'X0', 'X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X8',
      ...,
      'X375', 'X376', 'X377', 'X378', 'X379', 'X380', 'X382', 'X383', 'X38
4',
      'X385'],
      dtype='object', length=378)
```

In [51]:

df_train.info

Out[51]:

```
<bound method DataFrame.info of
8  ...  X375  X376  X377  X378  \
0      0  130.81  k  v  at  a  d  u  j  o  ...  0      0      1      0
1      6   88.53  k  t  av  e  d  y  l  o  ...  1      0      0      0
2      7   76.26  az  w  n  c  d  x  j  x  ...  0      0      0      0
3      9   80.62  az  t  n  f  d  x  l  e  ...  0      0      0      0
4     13   78.02  az  v  n  f  d  h  d  n  ...  0      0      0      0
...  ...  ...  ..  ..  ..  ..  ..  ..  ..  ..  ..  ...  ...  ...  ...
4204  8405  107.39  ak  s  as  c  d  aa  d  q  ...  1      0      0      0
4205  8406  108.77  j  o  t  d  d  aa  h  h  ...  0      1      0      0
4206  8412  109.22  ak  v  r  a  d  aa  g  e  ...  0      0      1      0
4207  8415   87.48  al  r  e  f  d  aa  l  u  ...  0      0      0      0
4208  8417  110.85  z  r  ae  c  d  aa  g  w  ...  1      0      0      0

      X379  X380  X382  X383  X384  X385
0         0      0      0      0      0      0
1         0      0      0      0      0      0
2         0      0      1      0      0      0
3         0      0      0      0      0      0
4         0      0      0      0      0      0
...  ...  ...  ...  ...  ...  ...
4204      0      0      0      0      0      0
4205      0      0      0      0      0      0
4206      0      0      0      0      0      0
4207      0      0      0      0      0      0
4208      0      0      0      0      0      0
```

[4209 rows x 378 columns]>

If for any column(s), the variance is equal to zero, then you need to remove those variable(s).

- Separating x_train and y_train dataset

In [52]:

```
xtrain=df_train.drop('y',axis=1)
ytrain=df_train['y']
```

In [53]:

```
xtrain.shape
```

Out[53]:

(4209, 377)

In [54]:

```
ytrain.shape
```

Out[54]:

(4209,)

In [55]:

```
xtrain.var()
```

Out[55]:

```
ID      5.941936e+06
X10     1.313092e-02
X11     0.000000e+00
X12     6.945713e-02
X13     5.462335e-02
...
X380    8.014579e-03
X382    7.546747e-03
X383    1.660732e-03
X384    4.750593e-04
X385    1.423823e-03
Length: 369, dtype: float64
```

In [56]:

```
df_test.head()
```

Out[56]:

	ID	X0	X1	X2	X3	X4	X5	X6	X8	X10	...	X375	X376	X377	X378	X379	X380	X382
0	1	az	v	n	f	d	t	a	w	0	...	0	0	0	1	0	0	0
1	2	t	b	ai	a	d	b	g	y	0	...	0	0	1	0	0	0	0
2	3	az	v	as	f	d	a	j	j	0	...	0	0	0	1	0	0	0
3	4	az	l	n	f	d	z	l	n	0	...	0	0	0	1	0	0	0
4	5	w	s	as	c	d	y	i	m	0	...	1	0	0	0	0	0	0

5 rows × 377 columns



In [57]:

```
df_test.shape
```

Out[57]:

(4209, 377)


In [58]:

```
df_test.describe()
```

Out[58]:

	ID	X10	X11	X12	X13	X14	
count	4209.000000	4209.000000	4209.000000	4209.000000	4209.000000	4209.000000	4209.000000
mean	4211.039202	0.019007	0.000238	0.074364	0.061060	0.427893	0.000000
std	2423.078926	0.136565	0.015414	0.262394	0.239468	0.494832	0.026000
min	1.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	2115.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
50%	4202.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
75%	6310.000000	0.000000	0.000000	0.000000	0.000000	1.000000	0.000000
max	8416.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000

8 rows × 369 columns



In [59]:

df_test.info

Out[59]:

```
<bound method DataFrame.info of
... X375 X376 X377 X378 \
0      1 az v n f d t a w 0 ... 0 0 0 1
1      2 t b ai a d b g y 0 ... 0 0 1 0
2      3 az v as f d a j j 0 ... 0 0 0 1
3      4 az l n f d z l n 0 ... 0 0 0 1
4      5 w s as c d y i m 0 ... 1 0 0 0
...    ... .. .. .. .. .. .. .. ..
4204 8410 aj h as f d aa j e 0 ... 0 0 0 0
4205 8411 t aa ai d d aa j y 0 ... 0 1 0 0
4206 8413 y v as f d aa d w 0 ... 0 0 0 0
4207 8414 ak v as a d aa c q 0 ... 0 0 1 0
4208 8416 t aa ai c d aa g r 0 ... 1 0 0 0

      X379 X380 X382 X383 X384 X385
0      0      0      0      0      0      0
1      0      0      0      0      0      0
2      0      0      0      0      0      0
3      0      0      0      0      0      0
4      0      0      0      0      0      0
...    ...    ...    ...    ...    ...    ...
4204      0      0      0      0      0      0
4205      0      0      0      0      0      0
4206      0      0      0      0      0      0
4207      0      0      0      0      0      0
4208      0      0      0      0      0      0
```

[4209 rows x 377 columns]>

In [60]:

df_test.columns

Out[60]:

```
Index(['ID', 'X0', 'X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X8', 'X10',
...
      'X375', 'X376', 'X377', 'X378', 'X379', 'X380', 'X382', 'X383', 'X38
4',
      'X385'],
      dtype='object', length=377)
```

In [61]:

```
for i in xtrain:
    if (xtrain[i].dtype=='O'):
        continue
    elif(xtrain.var()[i]==0):
        xtrain.drop(i,axis=1,inplace=True)
        df_test.drop(i,axis=1,inplace=True)
```

In [63]:

```
xtrain.head()
```

Out[63]:

	ID	X0	X1	X2	X3	X4	X5	X6	X8	X10	...	X375	X376	X377	X378	X379	X380	X382
0	0	k	v	at	a	d	u	j	o	0	...	0	0	1	0	0	0	0
1	6	k	t	av	e	d	y	l	o	0	...	1	0	0	0	0	0	0
2	7	az	w	n	c	d	x	j	x	0	...	0	0	0	0	0	0	1
3	9	az	t	n	f	d	x	l	e	0	...	0	0	0	0	0	0	0
4	13	az	v	n	f	d	h	d	n	0	...	0	0	0	0	0	0	0

5 rows × 365 columns



In [65]:

```
xtrain.shape
```

Out[65]:

(4209, 365)

Check for null and unique values for test and train sets.

In [68]:

```
xtrain.isnull().sum()
```

Out[68]:

```
ID      0
X0      0
X1      0
X2      0
X3      0
..
X380    0
X382    0
X383    0
X384    0
X385    0
Length: 365, dtype: int64
```

In [74]:

```
for i in xtrain:  
    if(xtrain[i].isnull().sum()!=0):  
        print(i)
```

ID
X0
X1
X2
X3
X4
X5
X6
X8
X10
X12
X13
X14
X15
X16
X17
X18
X19
X20
...

In [73]:

```
for i in df_test:  
    if(df_test[i].isnull().sum()!=0):  
        print(i)
```

In [75]:

```
ytrain.isnull().sum()
```

Out[75]:

0

In [76]:

```

for i in xtrain:
    if (xtrain[i].dtype=='O'):
        print(i)
        print(xtrain[i].unique())

```

X0

```

['k' 'az' 't' 'al' 'o' 'w' 'j' 'h' 's' 'n' 'ay' 'f' 'x' 'y' 'aj' 'ak' 'am'
 'z' 'q' 'at' 'ap' 'v' 'af' 'a' 'e' 'ai' 'd' 'aq' 'c' 'aa' 'ba' 'as' 'i'
 'r' 'b' 'ax' 'bc' 'u' 'ad' 'au' 'm' 'l' 'aw' 'ao' 'ac' 'g' 'ab']

```

X1

```

['v' 't' 'w' 'b' 'r' 'l' 's' 'aa' 'c' 'a' 'e' 'h' 'z' 'j' 'o' 'u' 'p' 'n'
 'i' 'y' 'd' 'f' 'm' 'k' 'g' 'q' 'ab']

```

X2

```

['at' 'av' 'n' 'e' 'as' 'aq' 'r' 'ai' 'ak' 'm' 'a' 'k' 'ae' 's' 'f' 'd'
 'ag' 'ay' 'ac' 'ap' 'g' 'i' 'aw' 'y' 'b' 'ao' 'al' 'h' 'x' 'au' 't' 'an'
 'z' 'ah' 'p' 'am' 'j' 'q' 'af' 'l' 'aa' 'c' 'o' 'ar']

```

X3

```

['a' 'e' 'c' 'f' 'd' 'b' 'g']

```

X4

```

['d' 'b' 'c' 'a']

```

X5

```

['u' 'y' 'x' 'h' 'g' 'f' 'j' 'i' 'd' 'c' 'af' 'ag' 'ab' 'ac' 'ad' 'ae'
 'ah' 'l' 'k' 'n' 'm' 'p' 'q' 's' 'r' 'v' 'w' 'o' 'aa']

```

X6

```

['j' 'l' 'd' 'h' 'i' 'a' 'g' 'c' 'k' 'e' 'f' 'b']

```

X8

```

['o' 'x' 'e' 'n' 's' 'a' 'h' 'p' 'm' 'k' 'd' 'i' 'v' 'j' 'b' 'q' 'w' 'g'
 'y' 'l' 'f' 'u' 'r' 't' 'c']

```

In [77]:

```
for i in df_test:
    if (df_test[i].dtype=='O'):
        print(i)
        print(df_test[i].unique())
```

```
X0
['az' 't' 'w' 'y' 'x' 'f' 'ap' 'o' 'ay' 'al' 'h' 'z' 'aj' 'd' 'v' 'ak'
 'ba' 'n' 'j' 's' 'af' 'ax' 'at' 'aq' 'av' 'm' 'k' 'a' 'e' 'ai' 'i' 'ag'
 'b' 'am' 'aw' 'as' 'r' 'ao' 'u' 'l' 'c' 'ad' 'au' 'bc' 'g' 'an' 'ae' 'p'
 'bb']
X1
['v' 'b' 'l' 's' 'aa' 'r' 'a' 'i' 'p' 'c' 'o' 'm' 'z' 'e' 'h' 'w' 'g' 'k'
 'y' 't' 'u' 'd' 'j' 'q' 'n' 'f' 'ab']
X2
['n' 'ai' 'as' 'ae' 's' 'b' 'e' 'ak' 'm' 'a' 'aq' 'ag' 'r' 'k' 'aj' 'ay'
 'ao' 'an' 'ac' 'af' 'ax' 'h' 'i' 'f' 'ap' 'p' 'au' 't' 'z' 'y' 'aw' 'd'
 'at' 'g' 'am' 'j' 'x' 'ab' 'w' 'q' 'ah' 'ad' 'al' 'av' 'u']
X3
['f' 'a' 'c' 'e' 'd' 'g' 'b']
X4
['d' 'b' 'a' 'c']
X5
['t' 'b' 'a' 'z' 'y' 'x' 'h' 'g' 'f' 'j' 'i' 'd' 'c' 'af' 'ag' 'ab' 'ac'
 'ad' 'ae' 'ah' 'l' 'k' 'n' 'm' 'p' 'q' 's' 'r' 'v' 'w' 'o' 'aa']
X6
['a' 'g' 'j' 'l' 'i' 'd' 'f' 'h' 'c' 'k' 'e' 'b']
X8
['w' 'y' 'j' 'n' 'm' 's' 'a' 'v' 'r' 'o' 't' 'h' 'c' 'k' 'p' 'u' 'd' 'g'
 'b' 'q' 'e' 'l' 'f' 'i' 'x']
```

Apply Label Encoder.

In [78]:

```
from sklearn import preprocessing
```

In [80]:

```
label_encoder = preprocessing.LabelEncoder()
```

In [81]:

```

for i in xtrain:
    if (xtrain[i].dtype!='int64'):
        xtrain[i]= label_encoder.fit_transform(xtrain[i])
    print(i)
    print(xtrain[i].unique())

```

X0

```

[32 20 40  9 36 43 31 29 39 35 19 27 44 45  7  8 10 46 37 15 12 42  5  0
 26  6 25 13 24  1 22 14 30 38 21 18 23 41  4 16 34 33 17 11  3 28  2]

```

X1

```

[23 21 24  3 19 13 20  1  4  0  6  9 26 11 16 22 17 15 10 25  5  7 14 12
  8 18  2]

```

X2

```

[17 19 34 25 16 14 38  7  8 33  0 31  3 39 26 24  5 21  2 13 27 29 20 42
 22 12  9 28 41 18 40 11 43  6 36 10 30 37  4 32  1 23 35 15]

```

X3

```

[0 4 2 5 3 1 6]

```

X4

```

[3 1 2 0]

```

X5

```

[24 28 27 12 11 10 14 13  9  8  5  6  1  2  3  4  7 16 15 18 17 20 21 23
 22 25 26 19  0]

```

X6

```

[ 9 11  3  7  8  0  6  2 10  4  5  1]

```

X8

```

[14 23  4 13 18  0  7 15 12 10  3  8 21  9  1 16 22  6 24 11  5 20 17 19
  2]

```

In [82]:

```

for i in df_test:
    if (df_test[i].dtype!='int64'):
        df_test[i]= label_encoder.fit_transform(df_test[i])
        print(i)
        print(df_test[i].unique())

```

X0
[21 42 45 47 46 29 12 38 20 8 31 48 6 27 44 7 23 37 33 41 3 19 15 13
17 36 34 0 28 5 32 4 22 9 18 14 40 11 43 35 26 1 16 25 30 10 2 39
24]

X1
[23 3 13 20 1 19 0 10 17 4 16 14 26 6 9 24 8 12 25 21 22 5 11 18
15 7 2]

X2
[34 8 17 4 38 24 26 10 33 0 16 6 37 32 9 23 14 13 2 5 22 29 30 27
15 35 19 39 44 43 21 25 18 28 12 31 42 1 41 36 7 3 11 20 40]

X3
[5 0 2 4 3 6 1]

X4
[3 1 0 2]

X5
[26 9 0 31 30 29 14 13 12 16 15 11 10 6 7 2 3 4 5 8 18 17 20 19
22 23 25 24 27 28 21 1]

X6
[0 6 9 11 8 3 5 7 2 10 4 1]

X8
[22 24 9 13 12 18 0 21 17 14 19 7 2 10 15 20 3 6 1 16 4 11 5 8
23]

Perform dimensionality reduction.

In [83]:

```
std_scalar=preprocessing.StandardScaler()
```

In [84]:

```
xtrain = std_scalar.fit_transform(xtrain)
```

In [85]:

```
df_test = std_scalar.fit_transform(df_test)
```

In [86]:

```
from sklearn.decomposition import PCA
```

In [87]:

```
pca=PCA()
```

In [88]:

```
xtrain_pca=pca.fit_transform(xtrain)
```

In [89]:

```
df_test_pca=pca.fit_transform(df_test)
```

In [90]:

```
pca.explained_variance_ratio_
```

Out[90]:

```
array([7.12486653e-02, 5.62415084e-02, 4.79979004e-02, 3.48818222e-02,
       3.26672655e-02, 3.17006619e-02, 2.82351002e-02, 2.14035131e-02,
       1.91298040e-02, 1.74214677e-02, 1.66286781e-02, 1.64393781e-02,
       1.49384640e-02, 1.38777649e-02, 1.35286203e-02, 1.27571090e-02,
       1.22387389e-02, 1.17690111e-02, 1.10581757e-02, 1.06911106e-02,
       1.02639043e-02, 9.37624352e-03, 9.29049134e-03, 9.05267345e-03,
       8.59958787e-03, 8.54160869e-03, 7.82044339e-03, 7.48747967e-03,
       7.34403625e-03, 7.23314434e-03, 7.03553099e-03, 6.87577048e-03,
       6.75483979e-03, 6.50479088e-03, 6.37950579e-03, 6.19292941e-03,
       6.13882865e-03, 6.03783025e-03, 5.99492525e-03, 5.89027248e-03,
       5.55576539e-03, 5.48572168e-03, 5.28121707e-03, 5.25910912e-03,
       5.12864286e-03, 5.02484447e-03, 4.94931135e-03, 4.76961699e-03,
       4.71278343e-03, 4.63198213e-03, 4.55830740e-03, 4.47537732e-03,
       4.35658769e-03, 4.28679476e-03, 4.22548636e-03, 4.04582798e-03,
       4.00033185e-03, 3.93454157e-03, 3.90301437e-03, 3.81967831e-03,
       3.80986083e-03, 3.72237302e-03, 3.65473153e-03, 3.64320550e-03,
       3.61870358e-03, 3.55727206e-03, 3.46128792e-03, 3.41047056e-03,
       3.37771387e-03, 3.36256219e-03, 3.31946796e-03, 3.30330864e-03,
       3.27652584e-03, 3.22861002e-03, 3.19479418e-03, 3.16074032e-03,
       3.10798244e-03, 3.10645053e-03, 3.07178635e-03, 3.05166619e-03,
       2.99650329e-03, 2.97911342e-03, 2.95233531e-03, 2.92968079e-03,
       2.88970794e-03, 2.88161868e-03, 2.85350724e-03, 2.83014050e-03,
       2.79960822e-03, 2.79022195e-03, 2.74931322e-03, 2.73849941e-03,
       2.69588542e-03, 2.68941004e-03, 2.64962483e-03, 2.62783267e-03,
       2.60106610e-03, 2.59075141e-03, 2.56129618e-03, 2.52025181e-03,
       2.48444007e-03, 2.46570355e-03, 2.45669962e-03, 2.43955390e-03,
       2.42671765e-03, 2.39406272e-03, 2.37753636e-03, 2.35029790e-03,
       2.32810902e-03, 2.26482158e-03, 2.24156399e-03, 2.23458252e-03,
       2.21686329e-03, 2.19543397e-03, 2.16864220e-03, 2.13939602e-03,
       2.10540805e-03, 2.09147504e-03, 2.08213708e-03, 2.05483315e-03,
       2.03818073e-03, 2.00300911e-03, 1.94897936e-03, 1.93704004e-03,
       1.93230270e-03, 1.92060136e-03, 1.88244702e-03, 1.83813991e-03,
       1.80528747e-03, 1.78224784e-03, 1.74610028e-03, 1.73211634e-03,
       1.71216167e-03, 1.67986342e-03, 1.64771357e-03, 1.61453168e-03,
       1.59628090e-03, 1.54925339e-03, 1.49466381e-03, 1.47998598e-03,
       1.47003925e-03, 1.43038582e-03, 1.41599576e-03, 1.38297063e-03,
       1.37771035e-03, 1.36436365e-03, 1.33269348e-03, 1.30417942e-03,
       1.27480706e-03, 1.26589598e-03, 1.24773234e-03, 1.23094905e-03,
       1.19234860e-03, 1.17073219e-03, 1.15206168e-03, 1.12350066e-03,
       1.10748818e-03, 1.10150579e-03, 1.06583699e-03, 1.04346695e-03,
       1.02463714e-03, 9.95758731e-04, 9.71998658e-04, 9.57689300e-04,
       9.30738127e-04, 9.16635380e-04, 8.99017230e-04, 8.76356461e-04,
       8.56126178e-04, 8.30582393e-04, 8.21694292e-04, 8.02493814e-04,
       7.76511877e-04, 7.74551191e-04, 7.43885041e-04, 7.28663565e-04,
       7.24732922e-04, 6.79961784e-04, 6.73459240e-04, 6.61756500e-04,
       6.39414138e-04, 6.25266526e-04, 6.03204345e-04, 5.79144335e-04,
       5.71735111e-04, 5.45853876e-04, 5.35603334e-04, 5.24151043e-04,
       5.14994605e-04, 4.97477026e-04, 4.89448595e-04, 4.76733679e-04,
       4.65189068e-04, 4.50214685e-04, 4.37565516e-04, 4.31458477e-04,
       4.23318421e-04, 4.08333193e-04, 3.97759383e-04, 3.80890485e-04,
       3.73848521e-04, 3.57616625e-04, 3.49990343e-04, 3.42519877e-04,
       3.24772837e-04, 3.18177219e-04, 3.13961066e-04, 3.09515606e-04,
       3.01351052e-04, 2.98079849e-04, 2.85629703e-04, 2.76391851e-04,
       2.64331204e-04, 2.47022641e-04, 2.40246893e-04, 2.25020442e-04,
       2.10275717e-04, 2.05191206e-04, 1.91552886e-04, 1.86367134e-04,
```

Downloaded from <http://ajph.org/> on November 10, 2015

In [92]:

```
pca=PCA(n_components=200)

xtrain_pca=pca.fit_transform(xtrain)

df_test_pca=pca.fit_transform(df_test)

pca.explained_variance_ratio_
```

Out[92]:

```
array([0.07124867, 0.05624151, 0.0479979 , 0.03488182, 0.03266727,
        0.03170066, 0.0282351 , 0.02140351, 0.0191298 , 0.01742147,
        0.01662868, 0.01643938, 0.01493846, 0.01387776, 0.01352862,
        0.01275711, 0.01223874, 0.01176901, 0.01105818, 0.01069111,
        0.0102639 , 0.00937624, 0.00929049, 0.00905267, 0.00859959,
        0.00854161, 0.00782044, 0.00748748, 0.00734404, 0.00723314,
        0.00703553, 0.00687577, 0.00675484, 0.00650479, 0.00637951,
        0.00619293, 0.00613883, 0.00603783, 0.00599493, 0.00589027,
        0.00555577, 0.00548572, 0.00528122, 0.00525911, 0.00512864,
        0.00502484, 0.00494931, 0.00476962, 0.00471278, 0.00463198,
        0.00455831, 0.00447538, 0.00435659, 0.00428679, 0.00422549,
        0.00404583, 0.00400033, 0.00393454, 0.00390301, 0.00381968,
        0.00380986, 0.00372237, 0.00365473, 0.00364321, 0.0036187 ,
        0.00355727, 0.00346129, 0.00341047, 0.00337771, 0.00336256,
        0.00331947, 0.00330331, 0.00327653, 0.00322861, 0.00319479,
        0.00316074, 0.00310798, 0.00310645, 0.00307179, 0.00305167,
        0.0029965 , 0.00297911, 0.00295234, 0.00292968, 0.00288971,
        0.00288162, 0.00285351, 0.00283014, 0.00279961, 0.00279022,
        0.00274931, 0.0027385 , 0.00269589, 0.00268941, 0.00264962,
        0.00262783, 0.00260107, 0.00259075, 0.0025613 , 0.00252025,
        0.00248444, 0.0024657 , 0.0024567 , 0.00243955, 0.00242672,
        0.00239406, 0.00237754, 0.0023503 , 0.00232811, 0.00226482,
        0.00224156, 0.00223458, 0.00221686, 0.00219543, 0.00216864,
        0.0021394 , 0.00210541, 0.00209147, 0.00208214, 0.00205483,
        0.00203818, 0.00200301, 0.00194898, 0.00193704, 0.0019323 ,
        0.0019206 , 0.00188245, 0.00183814, 0.00180529, 0.00178225,
        0.0017461 , 0.00173212, 0.00171216, 0.00167986, 0.00164771,
        0.00161453, 0.00159628, 0.00154925, 0.00149466, 0.00147999,
        0.00147004, 0.00143038, 0.001416 , 0.00138297, 0.00137771,
        0.00136436, 0.00133269, 0.00130418, 0.00127481, 0.00126589,
        0.00124772, 0.00123095, 0.00119234, 0.00117073, 0.00115206,
        0.0011235 , 0.00110748, 0.0011015 , 0.00106583, 0.00104346,
        0.00102463, 0.00099575, 0.00097199, 0.00095768, 0.00093071,
        0.00091662, 0.00089901, 0.00087633, 0.00085608, 0.00083057,
        0.00082165, 0.00080243, 0.00077647, 0.00077451, 0.00074386,
        0.0007283 , 0.00072465, 0.00067978, 0.00067339, 0.00066167,
        0.00063934, 0.00062479, 0.00060306, 0.00057847, 0.0005714 ,
        0.00054553, 0.0005354 , 0.00052402, 0.00051364, 0.00049689,
        0.00048843, 0.00047489, 0.00046366, 0.00044856, 0.00043099,
        0.00042752, 0.00041881, 0.00040684, 0.00039639, 0.00037901])
```

Predict your test_df values using XGBoost.

In [93]:

```
import xgboost as xgb
```

In [94]:

```
xgb_model=xgb.XGBRegressor(objective='reg:linear', n_estimators = 10)
```

In [95]:

```
xgb_model.fit(xtrain_pca,ytrain)
```

[21:40:45] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.5.1/src/objective/regression_obj.cu:188: reg:linear is now deprecated in favor of reg:squarederror.

Out[95]:

```
XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,
             colsample_bynode=1, colsample_bytree=1, enable_categorical=False,
             gamma=0, gpu_id=-1, importance_type=None,
             interaction_constraints='', learning_rate=0.300000012,
             max_delta_step=0, max_depth=6, min_child_weight=1, missing=nan,
             monotone_constraints='()', n_estimators=10, n_jobs=4,
             num_parallel_tree=1, objective='reg:linear', predictor='auto',
             random_state=0, reg_alpha=0, reg_lambda=1, scale_pos_weight=1,
             subsample=1, tree_method='exact', validate_parameters=1,
             verbosity=None)
```

In [96]:

```
xgb_model.predict(df_test_pca)
```

Out[96]:

```
array([ 78.506714,  90.84197 ,  80.794586, ...,  98.53473 , 103.381874,
        90.35204 ], dtype=float32)
```

In [98]:

```
from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score
```

In [100]:

```
ytrain_pred_xgb2=xgb_model.predict((xtrain_pca))
np.sqrt(mean_squared_error(ytrain ,ytrain_pred_xgb2))
```

Out[100]:

```
7.7529378889435625
```

In [102]:

```
r2_score(ytrain ,ytrain_pred_xgb2)
```

Out[102]:

0.6260274887750187

In [103]:

```
pd.DataFrame({"Actual Train Values" : ytrain , "Model Predictions" : ytrain_pred_xgb2})
```

Out[103]:

	Actual Train Values	Model Predictions
0	130.81	111.508102
1	88.53	89.199455
2	76.26	77.150223
3	80.62	81.635017
4	78.02	78.241142
...
4204	107.39	102.831642
4205	108.77	104.847939
4206	109.22	107.898170
4207	87.48	89.820564
4208	110.85	93.163147

4209 rows × 2 columns

In [104]:

```
#####
```

In []: